

Adversarial client

objective: make the model
misclassify the image

*same optimizer, start image,
perturbation ball, query budget*

Benign client

objective: e.g. restore the
image and keep its label

queries

Classifier

model f behind
a prediction API

queries

Stateful detector

per-client history;
alarm $\Rightarrow$ block

outputs

queries

query-only detectors use the queries;
output-aware detectors use both

